

Adding in our own transfer layer by hand with AlexNet

```
In [7]: import random
import tensorflow as tf
import numpy as np
import os
from scipy import ndimage
import matplotlib.pyplot as plt

%matplotlib inline
```

Load in our previous exported model

```
In [8]: graph = tf.Graph()
with graph.as_default():
    importer = tf.train.import_meta_graph('saved_models/alex_vars.meta')

sess = tf.Session(graph=graph)
importer.restore(sess, 'saved_models/alex_vars')
```

Get handle to second-to-last layer in pre-built model

```
In [9]: fc7_op = graph.get_operation_by_name('fc7/relu')
fc7 = fc7_op.outputs[0]
```

```
In [10]: fc7.get_shape()
```

```
Out[10]: TensorShape([Dimension(None), Dimension(4096)])
```

Create new layer, attached to fc7

```

In [11]: # Create new final layer
with graph.as_default():
    x = graph.get_operation_by_name('input').outputs[0]

    with tf.name_scope('transfer'):
        labels = tf.placeholder(tf.int32, [None])
        one_hot_labels = tf.one_hot(labels, 2)

        with tf.name_scope('cat_dog_final_layer'):
            weights = tf.Variable(tf.truncated_normal([4096, 2],
stddev=0.001),
                                name='final_weights')
            biases = tf.Variable(tf.zeros([2]), name='final_biases')
            logits = tf.nn.xw_plus_b(fc7, weights, biases, name='logits')

            prediction = tf.nn.softmax(logits, name='cat_dog_softmax')
            cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits, one_hot_labels)
            loss = tf.reduce_mean(cross_entropy, name='cat_dog_loss')

            global_step = tf.Variable(0, trainable=False, name='global_step')
            inc_step = global_step.assign_add(1)

            cat_dog_variables = [weights, biases]
            train = tf.train.GradientDescentOptimizer(0.01).minimize(loss, global_step=global_step,
                                                                    var_list=cat_dog_variables)

            with tf.name_scope('accuracy'):
                label_prediction = tf.argmax(prediction, 1, name='predicted_label')
                correct_prediction = tf.equal(label_prediction, tf.argmax(one_hot_labels, 1))
                accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

            init = tf.initialize_all_variables()

```

```

In [12]: sess = tf.Session(graph=graph)
sess.run(init)

```

Get our training file names

```
In [13]: cat_files = [  
    'data/dogs_and_cats/cats/' + f  
    for  
    f  
    in  
    os.listdir('data/dogs_and_cats/cats')  
]  
  
dog_files = [  
    'data/dogs_and_cats/dogs/' + f  
    for  
    f  
    in  
    os.listdir('data/dogs_and_cats/dogs')  
]  
  
all_files = cat_files + dog_files
```

Shuffle and split into training/validation

```
In [14]: random.shuffle(all_files)
```

```
In [15]: num_files = len(all_files)  
valid_percentage = 0.3  
split = int(num_files * valid_percentage)  
valid_data = all_files[:split]  
train_data = all_files[split:]
```

```
In [16]: print('Number of training images: {}'.format(len(train_data)))  
print('Number of validation images: {}'.format(len(valid_data)))
```

```
Number of training images: 17500  
Number of validation images: 7500
```

Create generator to give us batches of data

```
In [27]: from tensorflow.python.framework import graph_util  
from tensorflow.python.framework import tensor_shape  
from tensorflow.python.platform import gfile  
from tensorflow.python.util import compat
```

```
In [28]: flip_left_right = True  
         random_crop = 1  
         random_scale = 1  
         random_brightness = 1  
         num_channels = 3  
         height = 227  
         width = 227  
         pixel_depth = 255.0
```

In [34]:

```

import ntpath

def get_batch(batch_size, data, max_epochs, should_distort=False):
    distort_graph = tf.Graph()
    with distort_graph.as_default():
        """
        From https://github.com/tensorflow/tensorflow/blob/master/tensorflow/
        xamples/image_retraining/retrain.py
        """

        jpeg_name = tf.placeholder(tf.string, name='DistortJPGInput')
        jpeg_data = tf.read_file(jpeg_name)
        decoded_image = tf.image.decode_jpeg(jpeg_data, channels=3)
        resized_image = tf.image.resize_images(decoded_image, (height, width))
        decoded_image_as_float = tf.cast(decoded_image, dtype=tf.float32)
        decoded_image_4d = tf.expand_dims(decoded_image_as_float, 0)
        margin_scale = 1.0 + (random_crop / 100.0)
        resize_scale = 1.0 + (random_scale / 100.0)
        margin_scale_value = tf.constant(margin_scale)
        resize_scale_value = tf.random_uniform(tensor_shape.scalar(),
                                                minval=1.0,
                                                maxval=resize_scale)

        scale_value = tf.mul(margin_scale_value, resize_scale_value)
        precrop_width = tf.mul(scale_value, width)
        precrop_height = tf.mul(scale_value, width)
        precrop_shape = tf.pack([precrop_height, precrop_width])
        precrop_shape_as_int = tf.cast(precrop_shape, dtype=tf.int32)
        precropped_image = tf.image.resize_bilinear(decoded_image_4d,
                                                    precrop_shape_as_int)
        precropped_image_3d = tf.squeeze(precropped_image, squeeze_dims=[0])
        cropped_image = tf.random_crop(precropped_image_3d,
                                       [width, width,
                                       num_channels])

        if flip_left_right:
            flipped_image = tf.image.random_flip_left_right(cropped_image)
        else:
            flipped_image = cropped_image
        brightness_min = 1.0 - (random_brightness / 100.0)
        brightness_max = 1.0 + (random_brightness / 100.0)
        brightness_value = tf.random_uniform(tensor_shape.scalar(),
                                                minval=brightness_min,
                                                maxval=brightness_max)

        brightened_image = tf.mul(flipped_image, brightness_value)
        distort_result = tf.expand_dims(brightened_image, 0, name='DistortResu
lt')

    distort_sess = tf.Session(graph=distort_graph)

    epoch = 0
    idx = 0
    while epoch < max_epochs:
        batch = []
        labels = []
        for i in range(batch_size):
            if idx + i >= len(data):
                random.shuffle(data)
                epoch += 1
                idx = 0

```

```
image_path = data[idx + i].encode()
if should_distort:
    val = distort_sess.run(distort_result,
                           feed_dict={jpeg_name: image_path})
else:
    val = distort_sess.run(resized_image,
                           feed_dict={jpeg_name: image_path})
if b'dog' in ntpath.basename(image_path):
    labels.append(1)
else:
    labels.append(0)
batch.append(val)
idx += batch_size
yield batch, labels
```

In [35]: `sess.run(init)`

Quick save of our model to view later

In [36]: `writer = tf.train.SummaryWriter('tensorboard/alexnet_retrain', graph=graph)`
`writer.close()`

Train our model!

```
In [37]: for data_batch, label_batch in get_batch(32, train_data, 1, should_distort=True):
    data_batch = np.squeeze(data_batch)
    feed_dict = {x: data_batch, labels: label_batch}
    err, acc, step, _ = sess.run([loss, accuracy, inc_step, train],
                                feed_dict=feed_dict)

    if step % 50 == 0:
        print("Step: {} \t Accuracy: {} \t Error: {}".format(step, acc, err))
```

```
Step: 50      Accuracy: 0.9375      Error: 0.0700903013349
Step: 100     Accuracy: 0.875       Error: 0.175362020731
Step: 150     Accuracy: 0.9375     Error: 0.150483578444
Step: 200     Accuracy: 0.9375     Error: 0.149619147182
Step: 250     Accuracy: 0.90625    Error: 0.124697074294
Step: 300     Accuracy: 0.90625    Error: 0.12353708595
Step: 350     Accuracy: 0.96875    Error: 0.187275096774
Step: 400     Accuracy: 1.0      Error: 0.0302645843476
Step: 450     Accuracy: 0.96875    Error: 0.108887523413
Step: 500     Accuracy: 1.0      Error: 0.0338761284947
Step: 550     Accuracy: 0.96875    Error: 0.0903983265162
Step: 600     Accuracy: 1.0      Error: 0.0455834493041
Step: 650     Accuracy: 0.90625    Error: 0.178182154894
Step: 700     Accuracy: 1.0      Error: 0.0333553180099
Step: 750     Accuracy: 0.90625    Error: 0.149356365204
Step: 800     Accuracy: 0.90625    Error: 0.187900155783
Step: 850     Accuracy: 0.96875    Error: 0.0794009119272
Step: 900     Accuracy: 1.0      Error: 0.0737376660109
Step: 950     Accuracy: 0.96875    Error: 0.148608088493
Step: 1000    Accuracy: 0.9375     Error: 0.0959873497486
Step: 1050    Accuracy: 0.9375     Error: 0.279115825891
```

Validate

```
In [38]: def check_accuracy(valid_data):
    batch_size = 50
    num_correct = 0
    total = len(valid_data)
    i = 0
    for data_batch, label_batch in get_batch(batch_size, valid_data, 1):
        feed_dict = {x: data_batch, labels: label_batch}
        correct_guesses = sess.run(correct_prediction,
                                    feed_dict=feed_dict)
        num_correct += np.sum(correct_guesses)
        i += batch_size
        if i % (batch_size * 10) == 0:
            print('\tIntermediate accuracy: {}'.format((float(num_correct) / float(i))))
    acc = num_correct / float(total)
    print('\nAccuracy: {}'.format(acc))
```


In [39]: `check_accuracy(valid_data)`

```

Intermediate accuracy: 0.96
Intermediate accuracy: 0.944
Intermediate accuracy: 0.939333333333
Intermediate accuracy: 0.9445
Intermediate accuracy: 0.9448
Intermediate accuracy: 0.945333333333
Intermediate accuracy: 0.946285714286
Intermediate accuracy: 0.945
Intermediate accuracy: 0.945555555556
Intermediate accuracy: 0.9452
Intermediate accuracy: 0.944727272727
Intermediate accuracy: 0.945666666667
Intermediate accuracy: 0.945846153846
Intermediate accuracy: 0.947142857143
Intermediate accuracy: 0.946533333333

```

Accuracy: 0.953066666667

Once again, let's inspect for fun

In [40]: `def spot_check():`
 `filename = random.choice(valid_data)`
 `image = ndimage.imread(filename)`
 `feed_dict = {x: [image]}`
 `guess = sess.run(label_prediction, feed_dict=feed_dict)`
 `if guess[0] == 1:`
 `print('Guess: dog')`
 `else:`
 `print('Guess: cat')`
 `plt.imshow(image)`
 `plt.show()`

In [43]: `spot_check()`

Guess: cat



In []: