

## Using AlexNet to get Encoded Vectors for Image Retrieval

```
In [1]: import random
import tensorflow as tf
import numpy as np
import os
from scipy import ndimage
import matplotlib.pyplot as plt
from sklearn.neighbors import NearestNeighbors

%matplotlib inline
```

### Load in our previous exported model

```
In [2]: graph = tf.Graph()
with graph.as_default():
    importer = tf.train.import_meta_graph('saved_models/alex_vars.meta')

sess = tf.Session(graph=graph)
importer.restore(sess, 'saved_models/alex_vars')
```

### Get handle to second-to-last layer in pre-built model

```
In [17]: fc7_op = graph.get_operation_by_name('fc7/relu')
fc7 = fc7_op.outputs[0]
```

```
In [18]: fc7.get_shape()
```

```
Out[18]: TensorShape([Dimension(None), Dimension(4096)])
```

### Create new layer, attached to fc7

```

In [19]: # Create new final layer
with graph.as_default():
    x = graph.get_operation_by_name('input').outputs[0]

    with tf.name_scope('transfer'):
        labels = tf.placeholder(tf.int32, [None])
        one_hot_labels = tf.one_hot(labels, 2)

        with tf.name_scope('cat_dog_final_layer'):
            weights = tf.Variable(tf.truncated_normal([4096, 2], stddev=0.001),
                                  name='final_weights')
            biases = tf.Variable(tf.zeros([2]), name='final_biases')
            logits = tf.nn.xw_plus_b(fc7, weights, biases, name='logits')

        prediction = tf.nn.softmax(logits, name='cat_dog_softmax')
        cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits, one_hot_labels, name='cat_dog_cross_entropy')
        loss = tf.reduce_mean(cross_entropy, name='cat_dog_loss')

        global_step = tf.Variable(0, trainable=False, name='global_step')
        inc_step = global_step.assign_add(1)

        cat_dog_variables = [weights, biases]
        train = tf.train.GradientDescentOptimizer(0.01).minimize(loss, global_step=global_step, var_list=cat_dog_variables)

    with tf.name_scope('accuracy'):
        label_prediction = tf.argmax(prediction, 1, name='predicted_label')
        correct_prediction = tf.equal(label_prediction, tf.argmax(one_hot_labels, 1, name='correct_label'))
        accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))

    init = tf.initialize_all_variables()

```

```

In [20]: sess = tf.Session(graph=graph)
sess.run(init)

```

## Get data, as before

```
In [27]: cat_files = [  
    'data/fashion_data/blouse/' + f  
    for  
    f  
    in  
    os.listdir('data/fashion_data/blouse')  
]  
  
dog_files = [  
    'data/fashion_data/robe/' + f  
    for  
    f  
    in  
    os.listdir('data/fashion_data/robe')  
]  
  
all_files = cat_files + dog_files
```

## Shuffle and split into training/validation

```
In [29]: random.shuffle(all_files)
```

```
num_files = len(all_files) valid_percentage = 0.3 split = int(num_files * valid_percentage) valid_data  
= all_files[split:] train_data = all_files[:split]
```

```
In [31]: print('Number of training images: {}'.format(len(train_data)))  
print('Number of validation images: {}'.format(len(valid_data)))
```

```
Number of training images: 5570  
Number of validation images: 2387
```

## Create generator to give us batches of data

```
In [32]: from tensorflow.python.framework import graph_util  
from tensorflow.python.framework import tensor_shape  
from tensorflow.python.platform import gfile  
from tensorflow.python.util import compat
```

```
In [33]: flip_left_right = True  
         random_crop = 1  
         random_scale = 1  
         random_brightness = 1  
         num_channels = 3  
         height = 227  
         width = 227  
         pixel_depth = 255.0
```

In [34]:

```

import ntpath

def get_batch(batch_size, data, max_epochs, should_distort=False):
    distort_graph = tf.Graph()
    with distort_graph.as_default():
        """
        From https://github.com/tensorflow/tensorflow/blob/master/tensorflow/exan
        """
        jpeg_name = tf.placeholder(tf.string, name='DistortJPGInput')
        jpeg_data = tf.read_file(jpeg_name)
        decoded_image = tf.image.decode_jpeg(jpeg_data, channels=3)
        resized_image = tf.image.resize_images(decoded_image, (height, width))
        decoded_image_as_float = tf.cast(decoded_image, dtype=tf.float32)
        decoded_image_4d = tf.expand_dims(decoded_image_as_float, 0)
        margin_scale = 1.0 + (random_crop / 100.0)
        resize_scale = 1.0 + (random_scale / 100.0)
        margin_scale_value = tf.constant(margin_scale)
        resize_scale_value = tf.random_uniform(tensor_shape.scalar(),
                                                minval=1.0,
                                                maxval=resize_scale)
        scale_value = tf.mul(margin_scale_value, resize_scale_value)
        precrop_width = tf.mul(scale_value, width)
        precrop_height = tf.mul(scale_value, width)
        precrop_shape = tf.pack([precrop_height, precrop_width])
        precrop_shape_as_int = tf.cast(precrop_shape, dtype=tf.int32)
        precropped_image = tf.image.resize_bilinear(decoded_image_4d,
                                                    precrop_shape_as_int)
        precropped_image_3d = tf.squeeze(precropped_image, squeeze_dims=[0])
        cropped_image = tf.random_crop(precropped_image_3d,
                                       [width, width,
                                       num_channels])

        if flip_left_right:
            flipped_image = tf.image.random_flip_left_right(cropped_image)
        else:
            flipped_image = cropped_image
        brightness_min = 1.0 - (random_brightness / 100.0)
        brightness_max = 1.0 + (random_brightness / 100.0)
        brightness_value = tf.random_uniform(tensor_shape.scalar(),
                                                minval=brightness_min,
                                                maxval=brightness_max)
        brightened_image = tf.mul(flipped_image, brightness_value)
        distort_result = tf.expand_dims(brightened_image, 0, name='DistortResult')

    distort_sess = tf.Session(graph=distort_graph)

    epoch = 0
    idx = 0
    while epoch < max_epochs:
        batch = []
        labels = []
        for i in range(batch_size):
            if idx + i >= len(data):
                random.shuffle(data)
                epoch += 1
                idx = 0
            image_path = data[idx + i].encode()
            if should_distort:

```

```

        val = distort_sess.run(distort_result,
                                feed_dict={jpeg_name: image_path})
    else:
        val = distort_sess.run(resized_image,
                                feed_dict={jpeg_name: image_path})
    if b'dog' in ntpath.basename(image_path):
        labels.append(1)
    else:
        labels.append(0)
    batch.append(val)
    idx += batch_size
    yield batch, labels

```

In [35]: `sess.run(init)`

## Quick save of our model to view later

In [36]: `writer = tf.train.SummaryWriter('tensorboard/alexnet_retrain', graph=graph)`  
`writer.close()`

## Train our model!

In [37]: `for data_batch, label_batch in get_batch(32, train_data, 1, should_distort=True):`  
 `data_batch = np.squeeze(data_batch)`  
 `feed_dict = {x: data_batch, labels: label_batch}`  
 `err, acc, step, _ = sess.run([loss, accuracy, inc_step, train],`  
 `feed_dict=feed_dict)`  
 `if step % 50 == 0:`  
 `print("Step: {} \t Accuracy: {} \t Error: {}".format(step, acc, err))`

Step: 50	Accuracy: 1.0	Error: 1.4651019228e-05
Step: 100	Accuracy: 1.0	Error: 2.3005895855e-05
Step: 150	Accuracy: 1.0	Error: 5.11963007739e-05
Step: 200	Accuracy: 1.0	Error: 0.000133767549414
Step: 250	Accuracy: 1.0	Error: 0.000319783517625
Step: 300	Accuracy: 1.0	Error: 1.86829911399e-05
Step: 350	Accuracy: 1.0	Error: 0.000107844040031

## Validate

```
In [38]: def check_accuracy(valid_data):
    batch_size = 50
    num_correct = 0
    total = len(valid_data)
    i = 0
    for data_batch, label_batch in get_batch(batch_size, valid_data, 1):
        feed_dict = {x: data_batch, labels: label_batch}
        correct_guesses = sess.run(correct_prediction,
                                   feed_dict=feed_dict)
        num_correct += np.sum(correct_guesses)
        i += batch_size
        if i % (batch_size * 10) == 0:
            print('\tIntermediate accuracy: {}'.format((float(num_correct) / float(
acc = num_correct / float(total)
print('\nAccuracy: {}'.format(acc))
```

```
In [39]: check_accuracy(valid_data)
```

```
Intermediate accuracy: 1.0
Intermediate accuracy: 1.0
Intermediate accuracy: 1.0
Intermediate accuracy: 1.0
```

```
Accuracy: 1.00544616674
```

## Decide how many examples want for our nearest neighbors model

```
In [63]: num_images = 2000
    #neighbor_list = all_files[:num_images]
    neighbor_list = valid_data[:num_images]
```

## Create empty NumPy array to fill with encoded vectors

```
In [64]: extracted_features = np.ndarray((num_images, fc7.get_shape()[1]))
```

## Fill said NumPy array



```
In [66]: for i, filename in enumerate(neighbor_list):  
         image = ndimage.imread(filename)  
         features = sess.run(fc7, feed_dict={x: [image]})  
         extracted_features[i:i+1] = features  
         if i % 250 == 0:  
             print(i)
```

```
0  
250  
500  
750  
1000  
1250  
1500  
1750
```

```
In [67]: len(extracted_features)
```

```
Out[67]: 2000
```

## Create Nearest Neighbors model!

```
In [95]: nbrs = NearestNeighbors(n_neighbors=3, algorithm='ball_tree').fit(extracted_features)
```

```
In [96]: distances, indices = nbrs.kneighbors(extracted_features)
```



```
In [97]: indices
```

```
Out[97]: array([[ 0, 827,  70],  
               [ 1, 1901, 1809],  
               [ 2, 317, 184],  
               ...,  
               [1997, 1877, 949],  
               [1998, 275, 57],  
               [1999, 70, 1026]])
```


## Print out the three nearest neighbors

```
In [98]: def show_neighbors(idx, indices, filenames):  
         neighbors = indices[idx]  
         for i, neighbor in enumerate(neighbors):  
             image = ndimage.imread(filenames[neighbor])  
             plt.figure(i)  
             plt.imshow(image)  
         plt.show()
```

```
In [111]: show_neighbors(random.randint(300, len(extracted_features)), indices, neighbor_li
```





```
In [114]: show_neighbors(random.randint(3, len(extracted_features)), indices, neighbor_list
```






```
In [122]: show_neighbors(random.randint(100, len(extracted_features)), indices, neighbor_li
```





```
In [128]: show_neighbors(random.randint(7, len(extracted_features)), indices, neighbor_list
```



```
In [164]: show_neighbors(random.randint(11, len(extracted_features)), indices, neighbor_list)
```



```
In [137]: show_neighbors(random.randint(50, len(extracted_features)), indices, neighbor_list)
```







```
In [174]: show_neighbors(random.randint(75, len(extracted_features)), indices, neighbor_list)
```



```
In [163]: show_neighbors(random.randint(200, len(extracted_features)), indices, neighbor_li
```



```
In [ ]:
```

















